

Ryan Reece, Ph.D.

Machine learning engineer / data scientist / physicist

Mountain View, CA

ryan.reece@gmail.com

[rreece.github.io](https://github.com/rreece)

[linkedin.com/in/ryanreece](https://www.linkedin.com/in/ryanreece)

+1-817-888-7010

EXPERIENCE

Staff Machine Learning Engineer | Oct 2022 - present (1 yr 1 mo)

[Tenstorrent](#), Santa Clara, CA

- Worked with full-stack engineers to develop a cloud for push-button provisioning of Tenstorrent accelerators for AI applications with kubernetes
- Lead the containerization of the software stack for AI applications
- [Tenstorrent started offering inference as a service on various NLP tasks through EdenAI APIs](#), a project I lead technically
- Developed inference API demos for various language (BERT, t5, [falcon](#)) and vision ([yolo](#)) models
- Helped benchmark performance of many models on multiple generations of Tenstorrent's accelerators
- Contributed in meetings with leadership to reviewing the product readiness of the software

Machine Learning Engineer | Apr 2018 - Aug 2022 (4 yrs 4 mos)

[Cerebras Systems](#), Sunnyvale, CA

- Developed end-to-end model references in both [pytorch](#) and [tensorflow](#), including the input data pipeline for Cerebras Wafer-Scale Engines
- Trained benchmark models and did exploratory optimization of various models for computer vision ([ResNets](#)) and NLP ([GNMT](#), [Transformer](#), [Linformer](#), [BERT](#), [RoBERTa](#), [GPT-2](#)); explored impacts of using [mixed precision](#), bucketing by sequence length, activation sparsity
- Model references and data pipeline code delivered to customers in the Model Zoo with several examples and detailed documentation
- Helped develop a new normalization layer, [OnlineNorm](#), that uses streaming statistics to allow normalization of activations with small batch sizes [[NeurIPS 2019](#)]
- Directly engaged and supported customers (from both national labs and industry) in meetings and on-sites; including helping Pittsburgh Supercomputing Center (PSC) select among scientific proposals for using their dual CS-1 system: [Neocortex](#)
- Triaged, explored, and tested customer-shared models; represented customer requirements to compiler engineers; debugged model and data pipeline issues for customers
- Co-authored a blog about [Getting started with PyTorch BERT models on the Cerebras CS-2 System](#)

Artificial Intelligence Fellow | Jan 2018 - Mar 2018 (2 mos)

[Insight Data Science](#), Palo Alto, CA

- Learned about data science and machine learning applications in a variety of business domains
- Developed cloud-based hyperparameter optimization platform: [HYPR.AI](#), for automating the testing of many ML models using AWS/Paperspace in docker containerized jobs

Postdoctoral Research Fellow | Jul 2013 - Aug 2017 (4 yrs 2 mos)

[Santa Cruz Institute for Particle Physics](#), The University of California, Santa Cruz, and

The European Council for Nuclear Research ([CERN](#)), Geneva, Switzerland

- 11 years (postdoc and Ph.D.) as a member of the ATLAS experiment, a 3000+ person collaboration looking for new physics in high energy proton-proton collisions at the Large Hadron Collider (LHC)
- Long involvement in codebase of more than 10 million lines of C++ and almost as many lines of Python
- [Expert in petabyte data reduction](#) (ATLAS ~10 PB/year), world-wide grid computing, and [data visualization](#) as a user and primary supporter of our group's 200-CPU computing cluster, accumulated more than [350k CPU-hours](#)
- Lead analysis groups as "Editor" in different searches for signals of supersymmetry and exotic decays, contributed to 6 research publications, and defended their approval

- 2015-17, full-time support the operations of the [data acquisition system \(DAQ\)](#) and detector monitoring systems of the SCT (a tracking sub-detector in ATLAS)
- 2016-17, built more [expertise in machine learning techniques](#), deep learning frameworks using Keras to build CNNs for [particle classification](#), and another project using sklearn for anomaly detection by [clustering with Gaussian Mixture Models](#)

Graduate Researcher | Jun 2006 - Jul 2013 (7 yrs)

[The University of Pennsylvania](#), Philadelphia, PA, and

The European Council for Nuclear Research ([CERN](#)), Geneva, Switzerland

- First summers as a student with Penn (2006-08) at CERN participating in the [integration and commissioning of custom electronics](#) for the Transition Radiation Tracker (TRT), the outermost sub-detector of the ATLAS tracker
- 2009-12, throughout most of the running of the LHC, rotated the [on-call responsibility](#) for the TRT DAQ
- Ph.D. research with the data from ATLAS focused on the identification of decays of tau leptons and their use in searches for new physics, a [pattern recognition](#) problem to identify a type of particle
- 2009-10, was the lead developer of the cut-based tau identification used with the first ATLAS data
- 2010-12, helped develop advanced tau identification using [Boosted Decision Trees \(BDTs\)](#) which superseded the above
- Knack for developing data analysis frameworks: *e.g.* [pyframe](#) has been used by several analyses in ATLAS
- The ATLAS and CMS experiments at the LHC [discovered the long-sought-after Higgs boson](#), evidence of which was announced on July 4, 2012 [[Physics Letters B](#), [arxiv:1207.7214](#)]

EDUCATION

- **Ph.D. Experimental Particle Physics**, The University of Pennsylvania (Philadelphia, PA), Jun 2006 - Jul 2013 thesis: "[A search for new physics in high-mass ditau events in the ATLAS detector](#)"
- **B.S. Physics with Honors**, The University of Texas (Austin, TX), Aug 2003 - May 2006 thesis: "Late pulsing in the Hamamatsu R1408 PMT used in the Sudbury Neutrino Observatory"

PUBLICATIONS

- Chiley, V. *et al.* (2019). Online normalization for training neural networks. [NeurIPS 2019](#). [[arxiv:1905.05894](#)]
- Albertsson, K. *et al.* (2018). Machine learning in high energy physics community white paper. [[arxiv:1807.02876](#)]
- As a member of the ATLAS collaboration since June 1, 2008, I am an "author" of more than 800 publications ([google scholar](#), [inspire](#)), however, my list of selected publications is here: [rreece.github.io/publications](#), but in particular:
 1. Search for supersymmetry in a final state containing two photons and missing transverse momentum in $\sqrt{s} = 13$ TeV pp collisions at the LHC using the ATLAS detector. [European Physical Journal C](#), 76, 517 (2016). [[arxiv:1606.09150](#)]
 2. Identification and energy calibration of hadronically decaying tau leptons with the ATLAS experiment in pp collisions at $\sqrt{s} = 8$ TeV. [European Physical Journal C](#), 75, 303 (2015). [[arxiv:1412.7086](#)]
 3. A search for high-mass resonances decaying to $\tau^+\tau^-$ with the ATLAS detector. [Physics Letters B](#), 719, 242-260 (2013). [[arxiv:1210.6604](#)]
 4. Performance of the ATLAS detector using first collision data. [Journal of High Energy Physics](#), 9, 56 (2010). [[arxiv:1005.5254](#)]

SKILLS

- **General:** deep learning (NLP and CV), statistical analysis, data visualization, data-driven modeling, anomaly detection, neural network classifiers, boosted decision trees, petabyte data reduction, object-oriented design, polymorphic interfaces, writing technical reports, working independently and in groups, presenting my ideas, graduate level physics and mathematics
- **Programming languages (fluent):** [C/C++/STL](#) (21+ years), [Python](#) (17+ years);
(**experienced**): javascript, SQL; **Markup languages:** [L^AT_EX](#), Markdown, (x)html with css

- **ML / Data science software:** [pytorch](#), [tensorflow](#), [keras](#), [HuggingFace](#), matplotlib, numpy, scipy, scikit-learn, pandas, jupyter, AWS (EC2, S3), docker, singularity, ROOT, RooStats, TMVA
- **General software:** Linux, bash, git, svn, UML, QT, Mathematica
- **Hobbies:** poker, philosophy, cycling, climbing

Last updated: November 18, 2023